Koutilya PNVR

5^{th} year Ph.D. | Computer Vision | University of Maryland, College Park

https://koutilya-pnvr.github.io
linkedin.com/in/pnvrk github.com/koutilya-pnvr
+1 202 460 6957 koutilya@terpmail.umd.edu
4320 Rowalt Dr, Apt 202, College Park, MD-20740

I am a 5^{th} year Ph.D. student, advised by **Prof. David Jacobs** at the University of Maryland, College Park. I am broadly interested in computer vision. Some of the areas I am currently working on are: **Domain adaptation, in particular between synthetic and real datasets; Semi-Supervised Learning, Knowledge-Distillation for regression, GANs; Depth, and Normal estimation**; Inverse Rendering; Video classification; and Inpainting. I am looking to further explore new research areas in computer vision.

Education

Current	3.93	**Ph.D.** | Computer Vision, University of Maryland, College Park
2017	-	Deep Learning Nanodegree Foundation, Udacity
2017	3.9	**M.S.** | Electrical and Computer Engineering, University of Maryland, College Park
2015	8.22	**B.Tech.** | Electrical Engineering, Indian Institute of Technology, Delhi, India

Publications

- Koutilya PNVR, Hao Zhou, and David Jacobs."**SharinGAN : Combining Synthetic and Real data for Unsupervised Geometry Estimation.**" In Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition, 2020.
- Hanson Alex, Koutilya Pnvr, Sanjukta Krishnagopal, and Larry Davis."**Bidirectional Convolutional LSTM for the Detection of Violence in Videos.**" The European Conference on Computer Vision (ECCV) Workshops, 2018.

Internship Experience

Dec 2020 – May 2020 | **Project NEON, STAR Labs, Campbell, CA**

- Worked with team of researchers and engineers on various audio-visual and self-supervised learning techniques.
- Prototyped novel learning algorithms in large scale production system for various audio and video synthesis approaches.
- Integrated solutions in cross language technology stack consisting of Python, C++ and CUDA.

Audio-visual Computer Vision | Self-supervision for Audio-Video modalities | Deep Learning | Pytorch

Research Experience

Present – Aug. 2017 | **5^{th} year Ph.D. | Computer Vision, University of Maryland, College Park**

- Domain Adaptation between synthetic and real datasets for applications such as Monocular Depth Estimation of outdoor scenes and Face Normal Estimation.
- Violence detection in videos using a Bidirectional ConvLSTM network.
- Guided Inpainting using Generative Adversarial Networks that can enable the use of different car images as guides to edit cars in street view scenes.

Computer Vision | Deep Learning | Domain Adaptation | Depth and Normal Estimation | Semi-Supervised Learning | Knowledge-Distillation | GANs | Inverse Rendering | Video classification | Guided Inpainting

Dec. 2017 – Aug. 2016 | **Graduate Research Assistant | Geographical Sciences, University of Maryland, College Park**

- Crop yield prediction using EPIC model that utilizes data from the soil, site, remote sensing, daily and monthly weather for different regions in the US.
- Cropland data layer prediction based on history of crops using a Bidirectional ConvLSTM network.
- Development of a crop phenology algorithm that accurately predicts the timings of different stages of crop growth ranging from Greenup to Senescence.

Satellite Imagery | Remote Sensing | Crop Modeling | High Performance Computing | Crop Phenology

Skills

Programming Languages	**Python**, **MATLAB**, **R**, C
Deep Learning frameworks	**Pytorch**, **TensorFlow**, TFLearn, Keras
Image Libraries	**OpenCV**, **PIL**, **Scikit-Image**, **Matplotlib**, **Numpy**, **Lycon**
Other technical skills	**Tensorboard**, **OpenMPI**, **OpenMP**, **LaTeX**, **Linux**, **Visual Studio Code**, **Git**

Technical Projects

Domain Adaptation between synthetic and real datasets for Geometry Estimation

May 2019 - Dec 2019

- Implemented a novel way of utilizing GANs to train deep networks that can minimize the domain gap between synthetic and real images for performance improvement in geometry estimation tasks.
- Demonstrated the performance gain for Monocular Depth Estimation and Face Normal Estimation tasks.
- Reduced the absolute error in testsets of KITTI by 23.77% and Make3D by 6.45% over the state-of-the-art.
- Improved the face normal prediction by 4.3% for $Acc < 20$ deg metric for the Photoface dataset over the state-of-the-art.

Domain Adaptation | Monocular Depth Estimation | Face Normal Estimation | GANs | Synthetic and Real datasets | Pytorch | Tensorboard

Violence Detection in videos using Bidirectional ConvLSTM

April 2018 - June 2018

- Developed a novel Bidirectional Convolutional LSTM network followed by an elementwise max-pooling layer to obtain better Spatio-temporal representations for detecting violence in videos.
- Demonstrated the superiority of our method on the Hockey fights, Movies and Violent-Flows datasets over previous state-of-the-art methods.
- Signified the importance of all our modules : BiConvLSTM, elementwise maxpool, temporal encoding via ablation studies.

Violence Detection | Bidirectional ConvLSTM | Elementwise Max Pool | Spatio-temporal encoding | Pytorch

Pose Classification for Human Faces using Convolutional Neural Networks (CNNs)

May 2017 - May 2017

- Implemented a CNN similar to LeNet-5 in TensorFlow aimed for pose classification in human faces.
- The architecture includes Convolutions, Max pooling and Fully connected layers with the ReLU activation function.
- Adopted and studied the Dropout technique in the network to prevent overfitting and to generalize well to the testset.
- Initially validated the architecture on CIFAR-10 with a test accuracy of 74% and finally obtained a pose classification accuracy of 95.96% for CMU faces dataset.

CNNs | Classification | Deep Learning | TensorFlow | CIFAR | CMU Faces | LeNet-5

MNIST Data classification using CNNs and SVMs

May. 2017 - May. 2017

- Implemented a CNN similar to LeNet-5 in TensorFlow for classifying MNIST with a test accuracy close to 99%.
- Chosen the best model via performance on the validation set by tweaking learning rate and batch size hyperparameters.
- Studied the same by training various kernel SVMs and achieved a best performance of 94.39% using RBF kernel.

CNNs | Classification | Deep Learning | TensorFlow | MNIST | LeNet-5 | SVMs | RBF kernel

Parallel implementation of SMO and Model Selection Algorithms for SVMs

Mar 2017 - May 2017

- Efficiently parallelized the Sequential Minimal Optimization (SMO) and Model Selection (MS) algorithms that define an optimal SVM classifier using OpenMPI, OpenMP and Hybrid frameworks in C++ on deepthought2 cluster.
- Studied the parallel version on several cluster configurations such as nodes and threads/core and over several datasets, such as MNIST, not-MNIST, Cordrna, A9a, Splice, etc.
- Model selection is performed on 900 models and the best one is chosen that gives the least validation error.
- A huge speedup of 173 is observed for the Hybrid implementation of MS-SMO on not-MNIST for 16 node configuration.

OpenMP | OpenMPI | Sequential Minimal Optimization (SMO) | Support Vector Machines (SVM) | High Performance Computing

References

David Jacobs
Professor
Department of Computer Science and UMIACS
University of Maryland College Park
djacobs@umiacs.umd.edu
+1 301 405 0679

Hao Zhou
Ph.D., Computer Vision
Applied Scientist
AWS Rekognition
zhouho@amazon.com
https://zhhoper.github.io/

Varaprasad Bandaru
Associate Research Professor
Department of Geographical Sciences
University of Maryland College Park
vbandaru@umd.edu
+1 301 405 3074