Koutilya PNVR

Ph.D. Candidate | Computer Vision | University of Maryland College Park

https://koutilya-pnvr.github.io
linkedin.com/in/pnvrk github.com/koutilya-pnvr
+1 202 460 6957 koutilya@terpmail.umd.edu
4320 Rowalt Dr, Apt 202, College Park, MD-20740

I am a Ph.D. candidate, advised by **Prof. David Jacobs** at the University of Maryland, College Park. I am broadly interested in computer vision. I am currently working on : **Leveraging deep generative models such as GANs for geometry estimation (depth, normals) via unsupervised domain adaptation; latent diffusion models for text-based segmentation of real and AI generated images**. Some more areas of computer vision I am interested in are : Inverse Rendering; Video classification; and Inpainting. I am highly interested in collaborating with researchers from the industry to work on fascinating ideas in computer vision.

Education

Current	3.93	**Ph.D.** | Computer Vision, University of Maryland, College Park
2017	3.9	**M.S.** | Electrical and Computer Engineering, University of Maryland, College Park
2015	8.22	**B.Tech.** | Electrical Engineering, Indian Institute of Technology, Delhi, India

Publications

- Koutilya PNVR, Bharat Singh, Pallabi Ghosh, Behjat Siddiquie, David Jacobs. "**LD-ZNet : A Latent Diffusion Approach for Text-Based Image Segmentation.**" In Proceedings of the International Conference on Computer Vision (ICCV), 2023.
- Koutilya PNVR, Hao Zhou, and David Jacobs."**SharinGAN : Combining Synthetic and Real data for Unsupervised Geometry Estimation.**" In Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR), 2020.
- Hanson Alex*, Koutilya PNVR*, Sanjukta Krishnagopal, and Larry Davis."**Bidirectional Convolutional LSTM for the Detection of Violence in Videos.**" The European Conference on Computer Vision (ECCV) Workshops, 2018.

Internship Experience

Dec 2022 – May 2022 | **Amazon GO, Seattle, WA**

- Explored the utility of various text-to-image generative models for text-based segmentation of images.
- Demonstrated the use of latent diffusion models (LDMs) pretrained on large-scale internet data for text-based segmentation.
- Proposed novel ways to utilize features from the internal stages of the LDM to improve the segmentation performance by nearly **6% on real images and 20% on AI-generated images**.

Text-Based Image Segmentation | Real Images | AI generated Images | Diffusion models | Stable Diffusion

Dec 2020 – May 2020 | **Project NEON, STAR Labs, Campbell, CA**

- Worked with team of researchers on various audio-visual and self-supervised learning techniques.
- Prototyped novel learning algorithms in large scale production system for various audio and video synthesis approaches.
- Integrated solutions in cross language technology stack consisting of Python, C++ and CUDA.

Audio-visual Computer Vision | Self-supervision for Audio-Video modalities | Deep Learning | Pytorch

Research Experience

Present – Aug. 2017 | **Ph.D. Candidate | Computer Vision, University of Maryland, College Park**

- Text-based segmentation of real and AI generated images by leveraging text-to-image latent diffusion generative model pretrained on large-scale internet data.
- Self-training methods such as knowledge-distillation targeted for monocular depth estimation.
- Domain Adaptation between synthetic and real datasets for applications such as Monocular Depth Estimation of outdoor scenes and Face Normal Estimation.
- Violence detection in videos using a Bidirectional ConvLSTM network.
- Guided Inpainting using Generative Adversarial Networks that can enable the use of different car images as guides to edit cars in street view scenes.

Computer Vision | Generative models - GANs and Diffusion models | Text-Based Image Segmentation | Domain Adaptation | Depth and Normal Estimation | Knowledge-Distillation | Inverse Rendering | Video classification | Guided Inpainting

Skills

Deep Learning frameworks	**Pytorch**, **TensorFlow**, TFLearn, Keras
Other technical skills	**Tensorboard**, **OpenMPI**, **OpenMP**, **LaTeX**, **Linux**, **Visual Studio Code**, **Git**

Technical Projects

Domain Adaptation between synthetic and real datasets for Geometry Estimation

May 2019 - Dec 2019

- Implemented a novel way of utilizing GANs to train deep networks that can minimize the domain gap between synthetic and real images for performance improvement in geometry estimation tasks.
- Demonstrated the performance gain for Monocular Depth Estimation and Face Normal Estimation tasks.
- Reduced the absolute error in testsets of KITTI by 23.77% and Make3D by 6.45% over the state-of-the-art.
- Improved the face normal prediction by 4.3% for $Acc < 20$ deg metric for the Photoface dataset over the state-of-the-art.

Domain Adaptation | Monocular Depth Estimation | Face Normal Estimation | GANs | Synthetic and Real datasets | Pytorch | Tensorboard

Violence Detection in videos using Bidirectional ConvLSTM

April 2018 - June 2018

- Developed a novel Bidirectional Convolutional LSTM network followed by an elementwise max-pooling layer to obtain better Spatio-temporal representations for detecting violence in videos.
- Demonstrated the superiority of our method on the Hockey fights, Movies and Violent-Flows datasets over previous state-of-the-art methods.
- Signified the importance of all our modules : BiConvLSTM, elementwise maxpool, temporal encoding via ablation studies.

Violence Detection | Bidirectional ConvLSTM | Elementwise Max Pool | Spatio-temporal encoding | Pytorch

Parallel implementation of SMO and Model Selection Algorithms for SVMs

Mar 2017 - May 2017

- Efficiently parallelized the Sequential Minimal Optimization (SMO) and Model Selection (MS) algorithms that define an optimal SVM classifier using OpenMPI, OpenMP and Hybrid frameworks in C++ on deepthought2 cluster.
- Studied the parallel version on several cluster configurations such as nodes and threads/core and over several datasets, such as MNIST, not-MNIST, Cordrna, A9a, Splice, etc.
- Model selection is performed on 900 models and the best one is chosen that gives the least validation error.
- A huge speedup of 173 is observed for the Hybrid implementation of MS-SMO on not-MNIST for 16 node configuration.

OpenMP | OpenMPI | Sequential Minimal Optimization (SMO) | Support Vector Machines (SVM) | High Performance Computing

References

David Jacobs
Professor
Department of Computer Science and UMIACS
University of Maryland College Park
djacobs@umiacs.umd.edu
+1 301 405 0679

Hao Zhou
Ph.D., Computer Vision
Applied Scientist
AWS Rekognition
zhouho@amazon.com
https://zhhoper.github.io/

Varaprasad Bandaru
Associate Research Professor
Department of Geographical Sciences
University of Maryland College Park
vbandaru@umd.edu
+1 301 405 3074