KOUTILYA PNVR
Ph.D. Candidate | Computer Vision | University of Maryland College Park

🌐 https://koutilya-pnvr.github.io
in linkedin.com/in/pnvrk   ⊙ github.com/koutilya-pnvr
▫ +1 202 460 6957   @ koutilya@terpmail.umd.edu
📍 4320 Rowalt Dr, Apt 202, College Park, MD-20740

I am a Ph.D. candidate, advised by **Prof. David Jacobs** at the University of Maryland, College Park. I am broadly interested in computer vision. I am currently working on : **Leveraging deep generative models such as GANs for geometry estimation (depth, normals) via unsupervised domain adaptation ; latent diffusion models for text-based segmentation of real and AI generated images**. Some more areas of computer vision I am interested in are : Inverse Rendering ; Video classification ; and Inpainting. I am highly interested in collaborating with researchers from the industry to work on fascinating ideas in computer vision.

🎓 EDUCATION

Current	3.93	**Ph.D.**	Computer Vision, University of Maryland, College Park
2017	3.9	**M.S.**	Electrical and Computer Engineering, University of Maryland, College Park
2015	8.22	**B.Tech.**	Electrical Engineering, Indian Institute of Technology, Delhi, India

📄 PUBLICATIONS

> Koutilya PNVR, Bharat Singh, Pallabi Ghosh, Behjat Siddiquie, David Jacobs. "**LD-ZNet : A Latent Diffusion Approach for Text-Based Image Segmentation.**" arXiv preprint arXiv:2303.12343, 2023.

> Koutilya PNVR, Hao Zhou, and David Jacobs. "**SharinGAN : Combining Synthetic and Real data for Unsupervised Geometry Estimation.**" In Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition, 2020.

> Hanson Alex*, Koutilya PNVR*, Sanjukta Krishnagopal, and Larry Davis. "**Bidirectional Convolutional LSTM for the Detection of Violence in Videos.**" The European Conference on Computer Vision (ECCV) Workshops, 2018.

💼 INTERNSHIP EXPERIENCE

Dec 2022 May 2022	**Amazon GO, Seattle, WA** > Explored the utility of various text-to-image generative models for text-based segmentation of images. > Demonstrated the use of latent diffusion models (LDMs) pretrained on large-scale internet data for text-based segmentation. > Proposed novel ways to utilize features from the internal stages of the LDM to improve the segmentation performance by nearly **6% on real images and 20% on AI-generated images**. Text-Based Image Segmentation · Real Images · AI generated Images · Diffusion models · Stable Diffusion
Dec 2020 May 2020	**Project NEON, STAR LABS, Campbell, CA** > Worked with team of researchers on various audio-visual and self-supervised learning techniques. > Prototyped novel learning algorithms in large scale production system for various audio and video synthesis approaches. > Integrated solutions in cross language technology stack consisting of Python, C++ and CUDA. Audio-visual Computer Vision · Self-supervision for Audio-Video modalities · Deep Learning · Pytorch

💼 RESEARCH EXPERIENCE

Present Aug. 2017	**Ph.D. Candidate	Computer Vision, UNIVERSITY OF MARYLAND, College Park** > Text-based segmentation of real and AI generated images by leveraging text-to-image latent diffusion generative model pretrained on large-scale internet data. > Self-training methods such as knowledge-distillation targeted for monocular depth estimation. > Domain Adaptation between synthetic and real datasets for applications such as Monocular Depth Estimation of outdoor scenes and Face Normal Estimation. > Violence detection in videos using a Bidirectional ConvLSTM network. > Guided Inpainting using Generative Adversarial Networks that can enable the use of different car images as guides to edit cars in street view scenes. Computer Vision · Generative models - GANs and Diffusion models · Text-Based Image Segmentation · Domain Adaptation · Depth and Normal Estimation · Knowledge-Distillation · Inverse Rendering · Video classification · Guided Inpainting

☰ SKILLS

Deep Learning frameworks	Pytorch, TensorFlow, TFLearn, Keras
Other technical skills	Tensorboard, OpenMPI, OpenMP, LaTeX, Linux, Visual Studio Code, Git

🖥 Technical Projects

Domain Adaptation between synthetic and real datasets for Geometry Estimation
May 2019 - Dec 2019

> - Implemented a novel way of utilizing GANs to train deep networks that can minimize the domain gap between synthetic and real images for performance improvement in geometry estimation tasks.
> - Demonstrated the performance gain for Monocular Depth Estimation and Face Normal Estimation tasks.
> - Reduced the absolute error in testsets of KITTI by 23.77% and Make3D by 6.45% over the state-of-the-art.
> - Improved the face normal prediction by 4.3% for $Acc < 20$ deg metric for the Photoface dataset over the state-of-the-art.

Domain Adaptation | Monocular Depth Estimation | Face Normal Estimation | GANs | Synthetic and Real datasets | Pytorch | Tensorboard

Violence Detection in videos using Bidirectional ConvLSTM
April 2018 - June 2018

> - Developed a novel Bidirectional Convolutional LSTM network followed by an elementwise max-pooling layer to obtain better Spatio-temporal representations for detecting violence in videos.
> - Demonstrated the superiority of our method on the Hockey fights, Movies and Violent-Flows datasets over previous state-of-the-art methods.
> - Signified the importance of all our modules : BiConvLSTM, elementwise maxpool, temporal encoding via ablation studies.

Violence Detection | Bidirectional ConvLSTM | Elementwise Max Pool | Spatio-temporal encoding | Pytorch

Parallel implementation of SMO and Model Selection Algorithms for SVMs
Mar 2017 - May 2017

> - Efficiently parallelized the Sequential Minimal Optimization (SMO) and Model Selection (MS) algorithms that define an optimal SVM classifier using OpenMPI, OpenMP and Hybrid frameworks in C++ on deepthought2 cluster.
> - Studied the parallel version on several cluster configurations such as nodes and threads/core and over several datasets, such as MNIST, not-MNIST, Cordrna, A9a, Splice, etc.
> - Model selection is performed on 900 models and the best one is chosen that gives the least validation error.
> - A huge speedup of 173 is observed for the Hybrid implementation of MS-SMO on not-MNIST for 16 node configuration.

OpenMP | OpenMPI | Sequential Minimal Optimization (SMO) | Support Vector Machines (SVM) | High Performance Computing

❝ References

David Jacobs
Professor
Department of Computer Science and UMIACS
University of Maryland College Park
@ djacobs@umiacs.umd.edu
📞 +1 301 405 0679

Hao Zhou
Ph.D., Computer Vision
Applied Scientist
AWS Rekognition
@ zhouho@amazon.com
🌐 https://zhhoper.github.io/

Varaprasad Bandaru
Associate Research Professor
Department of Geographical Sciences
University of Maryland College Park
@ vbandaru@umd.edu
📞 +1 301 405 3074